

## COL380: Assignment-0

### Code Profiling

Ishita Chaudhary  
2019CS10360  
January 16, 2023

### Introduction

The aim is to learn how to profile code, identify and analyze the hotspots, and suggest and implement further changes to the code to improve its performance. We will use the tool *perf* to do the same.

### Running Perf

The following experiments were run on CSS cluster and analyzed via the perf profiling tool.

We vary the number of threads in the program and run perf stat for thread values 1,4,8, 12,.....,32 and plotted the values of the total time elapsed and cycles as a function of the number of threads. As we can see below, the data is obtained for 32 threads.

```
11 run: classify
12 ./classify rfile dfile 1009072 32 3
13

PROBLEMS OUTPUT TERMINAL PORTS DEBUG CONSOLE

cs1190360@css5:~/COL380/A0/A0$ perf stat make run
./classify rfile dfile 1009072 32 3
38.3998 ms
34.7954 ms
33.8372 ms
3 iterations of 1009072 items in 1001 ranges with 32 threads: Fastest took 33.8372 ms, Average was 35.6775 ms

Performance counter stats for 'make run':

      163.22 msec task-clock                #   0.967 CPUs utilized
           15      context-switches        #   0.092 K/sec
              1      cpu-migrations        #   0.006 K/sec
          9,362      page-faults           #   0.057 M/sec
    713,558,726      cycles                #   4.372 GHz
  1,837,660,817      instructions          #   2.58  insn per cycle
    545,065,870      branches              # 3339.382 M/sec
   10,481,097      branch-misses           #   1.92% of all branches

0.168821837 seconds time elapsed

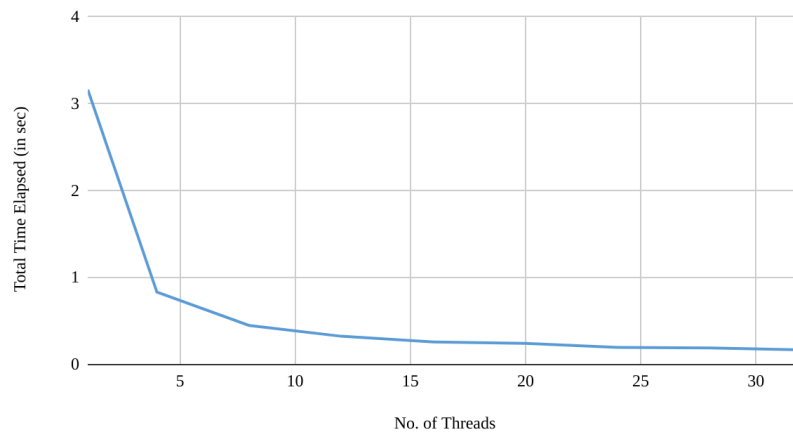
0.151995000 seconds user
0.012355000 seconds sys
```

The following values were obtained while running the experiment; the graphs of the data are attached below.


No. of Threads	Total Time Elapsed (in sec)	No. of Cycles
1	3.162526922	14,644,647,666
4	0.834439504	3,810,059,666
8	0.44995137	2,011,276,679
12	0.325487913	1,430,046,821

16	0.259278178	1,122,379,915
20	0.244099116	1,005,897,561
24	0.197652154	846,661,987
28	0.190470714	765,351,738
32	0.168821837	713,558,726

Total Time Elapsed (in sec) vs No. of Threads



No. of Cycles vs No. of Threads



As we increase the number of threads, the total time elapsed and the number of cycles decreases exponentially due to the SIMD structure of the core. Hence, concurrent execution of the same instruction in various threads can occur simultaneously.

For further experiments, we take four threads and ten repetitions.

Samples: 12K of event 'cycles', 4000 Hz, Event count (approx.): 14142682732	
classify /home/btech/cs1190360/COL380/A0/A0/classify [Percent: local period]	
Percent	
	lea -0x4(%r15,%rax,4),%r9
	lea 0x8(%rdx),%rax
	lea (%rax,%r14,1),%rdi
	→ jmp 355c <classify(Data&, Ranges const&, unsigned int)+0x2dc>
	nop
	add \$0x8,%rax
19.04	→ jne 3572 <classify(Data&, Ranges const&, unsigned int)+0x2f2>
	mov (%rdx),%rdx
	mov %esi,%r8d
	add (%r9),%r8d
	add \$0x1,%esi
0.01	mov %rdx,0x0(%r13,%r8,8)
	mov %rax,%rdx
2.75	→ jne 3558 <classify(Data&, Ranges const&, unsigned int)+0x2d8>
	add %r12d,%ecx
	cmp %ecx,%r10d
	→ jg 3538 <classify(Data&, Ranges const&, unsigned int)+0x2b8>
	mov -0x38(%rbp),%rbx
	xor %fs:0x28,%rbx
	mov -0x48(%rbp),%eax
	→ jne 360e <classify(Data&, Ranges const&, unsigned int)+0x38e>

Samples: 12K of event 'cycles', 4000 Hz, Event count (approx.): 14142682732	
classify /home/btech/cs1190360/COL380/A0/A0/classify [Percent: local period]	
Percent	
0.05	cltq
	lea (%r9,%rax,8),%rdi
0.01	mov 0x8(%r14),%eax
	mov (%rdi),%edx
	test %eax,%eax
0.01	→ jle 33f2 <classify(Data&, Ranges const&, unsigned int)+0x172>
	mov (%r14),%r11
	lea -0x1(%rax),%r15d
	xor %eax,%eax
	mov %eax,%ecx
8.35	cmp (%r11,%rax,8),%edx
11.46	→ jge 3390 <classify(Data&, Ranges const&, unsigned int)+0x110>
15.52	lea 0x1(%rax),%rcx
0.77	cmp %r15,%rax
	→ je 33f2 <classify(Data&, Ranges const&, unsigned int)+0x172>
	mov %rcx,%rax
11.13	→ jmp 33dc <classify(Data&, Ranges const&, unsigned int)+0x15c>
	mov -0x48(%rbp),%rax
	xor %ecx,%ecx
	→ jmp 339f <classify(Data&, Ranges const&, unsigned int)+0x11f>
	movabs \$0xffffffffffffffff,%rdx
	movslq 0x8(%r14),%rax
	cmp %rdx,%rax

Samples: 12K of event 'cycles', 4000 Hz, Event count (approx.): 14142682732 classify /home/btech/cs1190360/COL380/A0/A0/classify [Percent: local period]	
Percent	
	mov %eax,%r8d
	mov 0x8(%rbx),%r9
	mov %eax,%esi
	shl \$0x2,%r8
	→ jmp 33c3 <classify(Data&, Ranges const&, unsigned int)+0x143>
	nop
0.26	cmp 0x4(%r11,%rax,8),%edx
29.48	→ jg 33e4 <classify(Data&, Ranges const&, unsigned int)+0x164>
0.19	shl \$0x6,%rax
0.12	add -0x48(%rbp),%rax
	mov %ecx,0x4(%rdi)
0.26	mov (%rax),%rdx
0.12	cmp 0x8(%rax),%r10d
	→ jae 35ef <classify(Data&, Ranges const&, unsigned int)+0x36f>
0.01	lea (%rdx,%r8,1),%rax
	add %r12d,%esi
0.22	mov (%rax),%edx
0.02	add \$0x1,%edx
0.04	mov %edx,(%rax)
	mov %esi,%eax
0.01	cmp %esi,(%rbx)
	→ jbe 33fa <classify(Data&, Ranges const&, unsigned int)+0x17a>
	cltq

*Which assembly instruction takes the most CPU time?*

The "Jump if Greater" instruction (abbreviated as jg) takes the most CPU time.

*Can you map the instruction to the part of the source code it corresponds to?*

This corresponds to the jump within the for loop which maps each data point to its respective range.

```
#pragma omp parallel num_threads(numt)
{
    int tid = omp_get_thread_num(); // I am thread number tid
    for(int i=tid; i<D.ndata; i+=numt) // Threads together share-loop through all of Data
        int v = D.data[i].value = R.range(D.data[i].key); // For each data, find the interval of data's key,
                                                    // and store the interval id in value. D is changed.
    counts[v].increase(tid); // Found one key in interval v
}
```

We add the "-g" flag in the makefile along with CFLAGS to allow the perf report to show the source code and assembly instructions.

## Hotspot Analysis

*Report the top hotspot in your write-up (attach a screenshot of the perf report showing the code and the percentage of time taken).*

The top hotspot is the jump within the for loop which identifies the range of each data point, as reported above. The below screenshot confirms the above-mentioned claim, the address and offsets correspond to the for loop in the source code.

```

Samples: 12K of event 'cycles', 4000 Hz, Event count (approx.): 14386687119
classify /home/btech/cs1190360/COL380/A0/A0/classify [Percent: local period]
Percent  hi = b;
        }

        bool within(int val) const { // Return if val is within this range
        return(lo <= val && val <= hi);
0.27      cmp    0x4(%r11,%rax,8),%edx
28.14    → jg     33e4 <classify(Data&, Ranges const&, unsigned int)+0x164>
0.24      shl    $0x6,%rax
0.09      add    -0x48(%rbp),%rax
        _Z8classifyR4DataRK6Rangesj():
        mov     %ecx,0x4(%rdi)
        // and store the interval id in value. D is changed.
        counts[v].increase(tid); // Found one key in interval v
0.31      mov     (%rax),%rdx
        _ZN7Counter8increaseEj():
        assert(id < numcount);
0.10      cmp     0x8(%rax),%r10d
        → jae    35ef <classify(Data&, Ranges const&, unsigned int)+0x36f>
        _counts[id]++;
        lea     (%rdx,%r8,1),%rax
        _Z8classifyR4DataRK6Rangesj():
        for(int i=tid; i<D.ndata; i+=numt) { // Threads together share-loop through all of Data

```

```

Samples: 12K of event 'cycles', 4000 Hz, Event count (approx.): 14386687119
classify /home/btech/cs1190360/COL380/A0/A0/classify [Percent: local period]
Percent  mov     0x8(%rbx),%rdx
        if(D.data[d].value == r) // If the data item is in this interval
        D2.data[rangecount[r-1]+rcount++] = D.data[d]; // Copy it to the appropriate place in D2.
        movslq %ecx,%rax
        int rcount = 0;
        xor     %esi,%esi
        D2.data[rangecount[r-1]+rcount++] = D.data[d]; // Copy it to the appropriate place in D2.
        lea     -0x4(%r15,%rax,4),%r9
        lea     0x8(%rdx),%rax
        lea     (%rax,%r14,1),%rdi
        → jmp    355c <classify(Data&, Ranges const&, unsigned int)+0x2dc>
        nop
        add     $0x8,%rax
        if(D.data[d].value == r) // If the data item is in this interval
        cmp     %ecx,0x4(%rdx)
19.16    → jne    3572 <classify(Data&, Ranges const&, unsigned int)+0x2f2>
        D2.data[rangecount[r-1]+rcount++] = D.data[d]; // Copy it to the appropriate place in D2.
        mov     (%rdx),%rdx
        mov     %esi,%r8d
        add     (%r9),%r8d
        add     $0x1,%esi
        mov     %rdx,0x0(%r13,%r8,8)
        for(int d=0; d<D.ndata; d++) // For each interval, thread loops through all of data and

```

```

Samples: 12K of event 'cycles', 4000 Hz, Event count (approx.): 14386687119
classify /home/btech/cs1190360/COL380/A0/A0/classify [Percent: local period]
Percent  mov     (%r14),%r11
        lea     -0x1(%rax),%r15d
        xor     %eax,%eax
        mov     %eax,%ecx
        _ZNK5Range6withinEi():
        return(lo <= val && val <= hi);
8.54     cmp     (%r11,%rax,8),%edx
11.83    → jge    3390 <classify(Data&, Ranges const&, unsigned int)+0x110>
        _ZNK6Ranges5rangeEib():
        for(int r=0; r< num; r++) // Look through all intervals
15.00    lea     0x1(%rax),%rcx
0.92     cmp     %r15,%rax
        → je     33f2 <classify(Data&, Ranges const&, unsigned int)+0x172>
        mov     %rcx,%rax
12.19    → jmp    33dc <classify(Data&, Ranges const&, unsigned int)+0x15c>
        mov     -0x48(%rbp),%rax
        return r;
        }
        return BADRANGE; // Did not find any range
        xor     %ecx,%ecx
        → jmp    339f <classify(Data&, Ranges const&, unsigned int)+0x11f>
        _Z8classifyR4DataRK6Rangesj():
        }

```

*What is the prospective problem which makes this code snippet the top hotspot?  
Can the code be optimized to improve the performance of this hotspot? If it can be optimized, suggest the optimizations in the write-up.*

Yes, we can optimize the code to improve the performance of this hotspot. Instead of giving consequent data reads and writes to different threads, we should allocate work on contiguous blocks of data to each thread so that the cache miss rate while reading reduces significantly. This ensures that time is not wasted in unnecessarily loading and storing cache lines by different threads, thus saving time in each iteration.

After running the perf record with the desired flags, we get the following stats on the original code:

```
Available samples
10K branches
8K branch-misses
1K cache-misses
22 page-faults
10K cpu-cycles
```

### **Memory Profiling**

*Run perf mem record on the given code, and analyze the report generated.*

We get 437 CPU memory loads and ~1000 CPU memory stores. The report for each of them is attached below.

```
Available samples
437 cpu/mem-loads,ldlat=30/P
1K cpu/mem-stores/P
```

Overhead	Command	Shared Object	Symbol
49.48%	classify	libstdc++.so.6.0.28	[.] std::num_get<char, std::istreambuf_iterator<char, std::char_traits<char> >>::
14.91%	classify	classify	[.] classify
9.40%	classify	classify	[.] repeatrun
6.26%	classify	libstdc++.so.6.0.28	[.] std::istream::sentry::sentry
2.59%	classify	libstdc++.so.6.0.28	[.] std::istream::operator>>
1.49%	classify	classify	[.] readRanges
1.19%	classify	classify	[.] readData
1.12%	classify	libstdc++.so.6.0.28	[.] 0x000000000126e54
1.10%	classify	libstdc++.so.6.0.28	[.] 0x000000000125254
1.03%	classify	libstdc++.so.6.0.28	[.] 0x000000000125258
0.76%	classify	libc-2.31.so	[.] _int_malloc
0.68%	classify	[unknown]	[k] 0xfffffffff814a3868
0.59%	classify	[unknown]	[k] 0xfffffffff8147230f
0.59%	classify	libc-2.31.so	[.] malloc
0.55%	classify	[unknown]	[k] 0xfffffffff8148cbcb
0.52%	classify	[unknown]	[k] 0xfffffffff814b0dd3
0.51%	classify	[unknown]	[k] 0xfffffffff8145ce5f
0.49%	classify	[unknown]	[k] 0xfffffffff81284f01
0.48%	classify	[unknown]	[k] 0xfffffffff814adc36
0.46%	classify	[unknown]	[k] 0xfffffffff8148b989
0.45%	classify	[unknown]	[k] 0xfffffffff8148ba8c
0.38%	classify	[unknown]	[k] 0xfffffffff814b224d
0.37%	classify	[unknown]	[k] 0xfffffffff81708e96

Overhead	Command	Shared Object	Symbol
67.00%	classify	classify	[.] classify
11.38%	classify	libstdc++.so.6.0.28	[.] std::num_get<char, std::istreambuf_iterator<char, std::char_traits<char> >>::
5.10%	classify	[unknown]	[k] 0xfffffffff8145e27b
3.30%	classify	[unknown]	[k] 0xfffffffff814f7670
2.11%	classify	[unknown]	[k] 0xfffffffff814ae412
1.94%	classify	[unknown]	[k] 0xfffffffff8145ce5f
1.75%	classify	[unknown]	[k] 0xfffffffff814af18d
1.42%	classify	[unknown]	[k] 0xfffffffff8145e27f
1.05%	classify	[unknown]	[k] 0xfffffffff818023ee
0.50%	classify	[unknown]	[k] 0xfffffffff8133f5be
0.38%	classify	[unknown]	[k] 0xfffffffff8134231b
0.28%	classify	ld-2.31.so	[.] _dl_lookup_symbol_x
0.26%	classify	[unknown]	[k] 0xfffffffff814adae1
0.25%	classify	[unknown]	[k] 0xfffffffff814a38eb
0.25%	classify	[unknown]	[k] 0xfffffffff814f774e
0.25%	classify	[unknown]	[k] 0xfffffffff814f78c9
0.23%	classify	[unknown]	[k] 0xfffffffff814b109e
0.21%	classify	[unknown]	[k] 0xfffffffff814a37f4
0.16%	classify	[unknown]	[k] 0xfffffffff814f77ea
0.14%	make	[unknown]	[k] 0xfffffffff8129dbc9
0.13%	make	[unknown]	[k] 0xfffffffff814ae447
0.12%	classify	[unknown]	[k] 0xfffffffff814f7806
0.11%	make	[unknown]	[k] 0xfffffffff814f7890

We can observe that 67% of the memory loads are done in the function classify.

Report the top 2 hotspots (attach a screenshot of the perf report showing the code and percentage of time taken).

Percent	Code
28.14	→ jg 33e4 <classify(Data&, Ranges const&, unsigned int)+0x164>
0.24	shl \$0x6,%rax
0.09	add -0x48(%rbp),%rax
0.31	_Z8classifyR4DataRK6Rangesj(): mov %ecx,0x4(%rdi) // and store the interval id in value. D is changed. counts[v].increase(tid); // Found one key in interval v mov (%rax),%rdx
0.10	→ jae 35ef <classify(Data&, Ranges const&, unsigned int)+0x36f>
	_counts[id]++; lea (%rdx,%r8,1),%rax _Z8classifyR4DataRK6Rangesj(): for(int i=tid; i<D.ndata; i+=numt) { // Threads together share-loop through all of Data

```

Samples: 12K of event 'cycles', 4000 Hz, Event count (approx.): 14386687119
classify /home/btech/cs1190360/COL380/A0/A0/classify [Percent: local period]
Percent
    mov     0x8(%rbx),%rdx
    if(D.data[d].value == r) // If the data item is in this interval
    D2.data[rangeount[r-1]+rcount++] = D.data[d]; // Copy it to the appropriate place in D2.
    movslq  %ecx,%rax
    int rcount = 0;
    xor     %esi,%esi
    D2.data[rangeount[r-1]+rcount++] = D.data[d]; // Copy it to the appropriate place in D2.
    lea     -0x4(%r15,%rax,4),%r9
    lea     0x8(%rdx),%rax
    lea     (%rax,%r14,1),%rdi
    → jmp    355c <classify(Data&, Ranges const&, unsigned int)+0x2dc>
    nop
    add     $0x8,%rax
    if(D.data[d].value == r) // If the data item is in this interval
    cmp     %ecx,0x4(%rdx)
    19.16 → jne    3572 <classify(Data&, Ranges const&, unsigned int)+0x2f2>
    D2.data[rangeount[r-1]+rcount++] = D.data[d]; // Copy it to the appropriate place in D2.
    mov     (%rdx),%rdx
    mov     %esi,%r8d
    add     (%r9),%r8d
    add     $0x1,%esi
    mov     %rdx,0x0(%r13,%r8,8)
    for(int d=0; d<D.ndata; d++) // For each interval, thread loops through all of data and
    Press 'h' for help on key bindings

```

These are the two hotspots, and both of them are a part of the original classify function. The first one (which consumes 28.14% of execution time) is a for loop which finds the accurate data range for each data point and the second one is also a part of the for loop where we are finally writing the data which we return.

*Based on the hotspots you obtained, identify at least two issues in the code that makes it cache unfriendly and suggest improvements.*

- For each data, we find the interval of data's key, and store the interval id in value and then increase the count of the interval by one. While performing this, the threads together share-loop through all of the data in such a way that consequent reads occur from different cache lines resulting in a lot of cache-misses and subsequently un-necessary loads and stores. Instead we can allocate work on contiguous blocks of data to each thread so that the cache miss rate reduces while performing this task.
- On analysing the code, it was observed that while computing  $D_2$ , using 2 for loops is redundant and can be computed using a single loop improving time complexity from  $O(R \cdot D)$  to  $O(D)$

Cache Line Width is 64 bytes, int size is 4bytes and Item (struct) size is 8bytes (key (4bytes), value(4bytes)), that is, it is desirable to access 8 contiguous (64/8) Item values and 16 contiguous (64/4) int values. So, each thread iterates over 8 contiguous elements (item struct) to compute the range for each element (first for loop), and iterates over 16 contiguous elements (int) to compute the elements in a particular range.

For this, two variable Items Accessed By One Thread (=8) and, Ints Accessed By One Thread (=16) are used.

So, a thread will access all elements in L1 data cache and will then move to another row.



Run *perf mem* record after the improvements and submit the screenshots.

There were 437 CPU memory loads which were reduced to 412 CPU memory loads after optimizations. That is, the optimized code is better in terms of the number of times memory loads were done by the CPU.

```
Available samples
412 cpu/mem-loads,ldlat=30/P
1K cpu/mem-stores/P
```

After Optimization

```
Available samples
437 cpu/mem-loads,ldlat=30/P
1K cpu/mem-stores/P
```

Before Optimization

Run *perf* with the *cache misses* flag on the original code and the final code you obtain. Do you see an improvement? If not, suggest ways to further improve cache hit rate.

```
Available samples
1K cache-misses
```

```
Samples: 1K of event 'cache-misses', Event count (approx.): 3638446
Overhead Command Shared Object Symbol
31.06% classify classify [.] classify
17.24% classify classify [.] repeatrun
4.62% classify [unknown] [k] 0xffffffff81df90c0
3.92% classify [unknown] [k] 0xffffffff814fa233
3.61% classify [unknown] [k] 0xffffffff814f7880
3.13% classify [unknown] [k] 0xffffffff818023ee
2.40% classify [unknown] [k] 0xffffffff814fa219
2.14% classify [unknown] [k] 0xffffffff814fa285
1.57% classify [unknown] [k] 0xffffffff814f4894
1.35% classify [unknown] [k] 0xffffffff8148babd
1.24% classify [unknown] [k] 0xffffffff814f489f
1.22% classify [unknown] [k] 0xffffffff814fa2cd
1.12% classify [unknown] [k] 0xffffffff8148285d
1.11% classify [unknown] [k] 0xffffffff814fcc56
1.11% classify [unknown] [k] 0xffffffff814fa210
1.08% classify [unknown] [k] 0xffffffff814fcb80
```

Original Code

Samples: 1K of event 'cache-misses', Event count (approx.): 2827022

Overhead	Command	Shared Object	Symbol
21.11%	classify	classify	[.] classify
9.22%	classify	[unknown]	[k] 0xffffffff81df90c0
9.08%	classify	classify	[.] repeatrun
6.20%	classify	[unknown]	[k] 0xffffffff814fa219
4.32%	classify	[unknown]	[k] 0xffffffff814fa285
3.26%	classify	[unknown]	[k] 0xffffffff812c6c00
2.87%	classify	[unknown]	[k] 0xffffffff814f4880
2.83%	classify	[unknown]	[k] 0xffffffff8148bb17
2.77%	classify	[unknown]	[k] 0xffffffff812c6c18
2.73%	classify	[unknown]	[k] 0xffffffff818023ee
2.60%	classify	[unknown]	[k] 0xffffffff814f7924
2.22%	classify	[unknown]	[k] 0xffffffff814fcab9
2.19%	classify	[unknown]	[k] 0xffffffff814f780e
1.89%	classify	[unknown]	[k] 0xffffffff814fa210
1.80%	classify	[unknown]	[k] 0xffffffff814fa233
1.70%	classify	[unknown]	[k] 0xffffffff814f7910
1.51%	classify	[unknown]	[k] 0xffffffff814f48bd
1.45%	classify	[unknown]	[k] 0xffffffff814be8d6

### Optimized Code

Perf was used to obtain the cache misses in the original code as well as in the modified code. The initial code was cache-unfriendly due to instances of false-sharing. Instead of giving consequent data reads and writes to different threads, in the modified version, we allocate work on contiguous blocks of data to each thread so that the cache miss rate while reading reduces significantly.

The above screenshots verify that the percentage of cache misses in the classify function reduces in the modified code.

To **further improve** cache-hit rate, we must use perf with more fine-grained flags such as those in the screenshot attached below to analyze the cache performance at different levels(L1, L2, L3) and optimize the code accordingly, also keeping in mind the kind of architecture or hardware being used.

L1-dcache-load-misses	[Hardware cache event]
L1-dcache-loads	[Hardware cache event]
L1-dcache-stores	[Hardware cache event]
L1-icache-load-misses	[Hardware cache event]
LLC-load-misses	[Hardware cache event]
LLC-loads	[Hardware cache event]
LLC-store-misses	[Hardware cache event]
LLC-stores	[Hardware cache event]
branch-load-misses	[Hardware cache event]
branch-loads	[Hardware cache event]
dTLB-load-misses	[Hardware cache event]
dTLB-loads	[Hardware cache event]
dTLB-store-misses	[Hardware cache event]
dTLB-stores	[Hardware cache event]
iTLB-load-misses	[Hardware cache event]
node-load-misses	[Hardware cache event]
node-loads	[Hardware cache event]
node-store-misses	[Hardware cache event]
node-stores	[Hardware cache event]